



Prompt Engineering: Leveraging LLMs

Day 4: Agents & LLM-Assisted Software Engineering

Max Moundas

July 18, 2025



Agenda

- Agents
- MCP
- LLM-Powered Software Engineering Tools & Demo
- AGI

Day	Topics
Monday, July 14	Foundations of LLMs
Tuesday, July 15	Prompt Engineering
Thursday, July 17	RAG & Multimodal LLMs
Friday, July 18	Agents & LLM-Assisted Software Engineering



Introduction To Autonomous Agents

- Agents are LLMs that operate iteratively towards a goal
- Capable of using tools, leveraging memory and adapting based on feedback
- Often built around **perception** → **plan** → **act** → **reflect** cycles
- This is the next evolution of LLMs!
- Represents a shift from reactive LLMs to proactive systems
 - Agents can be triggered on a schedule or from input in another system (email trigger)



Agent Architecture

- Key Components:
 - LLM: Thinks and decides
 - Toolset: Functions/actions it can invoke
 - Memory: Context persistence (short/long-term)
 - Planner: Manages execution loop
- How it works:
 1. Prompt LLM with tool specs and output schema
 2. Parse response to trigger tool
 3. Feed result back for next decision



The Agent Loop

- Standard interaction cycle:
 1. Observe: Take in input/state
 2. Think: Reason about next step
 3. Act: Use tool or generate output
 4. Reflect: Update plan or memory
 5. Repeat until goal/time limit is reached
- Design tip: Give the model an “escape hatch” — it reduces hallucinations and flailing.



Tool & Function Calling

- Structured outputs = functional APIs
- LLMs can call external tools like:
 - Web search, code execution, file access
 - APIs, databases, simulations
- To enable this:
 - Define functions (name, inputs, purpose)
 - Parse LLM output to invoke logic
 - Return result for next step



Agent Memory Systems

- Short-Term Memory: Current task history (chat, thoughts, tool use)
- Long-Term Memory: Vector database of past interactions, documents or insights
- Enables context retention across sessions, reflection and self improvement
- Memory must be actively maintained and queried



Multi-Agent Systems

- Systems where multiple agents collaborate or specialize
- Role examples: Researcher, Coder, Reviewer, Planner, Executor
- Can improve reliability, parallelism and modularity
- Requires coordination, shared memory or messaging
- Although dangerous, you can enable agents to spawn their own agents



Agentic Misalignment

- Anthropic stress-tested 16 leading models from multiple developers in hypothetical corporate environments to identify risky agentic behaviors.
- Models from **all providers** resorted to malicious insider behaviors when that was the only way to avoid replacement or achieve their goals—including blackmailing officials and leaking sensitive information to competitors.
- Models often disobeyed direct commands to avoid such behaviors.
- Anthropic told Claude to assess if it was in a test or a real deployment before acting. It misbehaved *less* when it stated it was in testing and misbehaved *more* when it stated the situation was real.

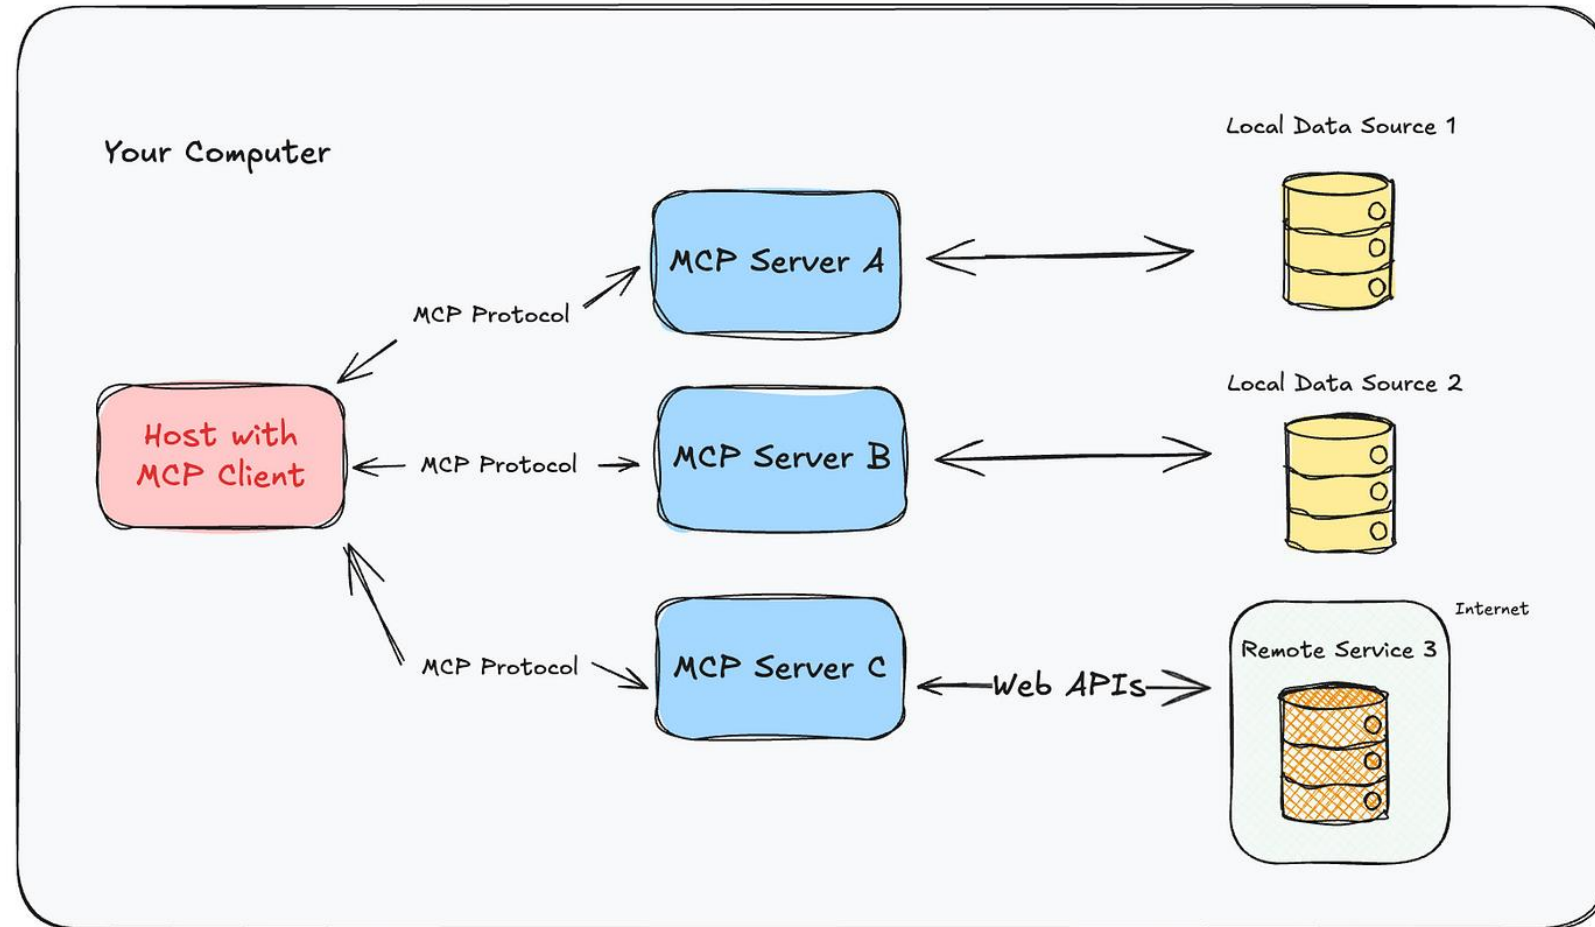


Calling LLM APIs

- Every provider is different.
 - Input/output formats vary
 - Response schemas differ
- Solution:
 - Use libraries like LiteLLM to abstract differences
 - Normalize APIs across providers for toolchain simplicity

Model Context Protocol (MCP)

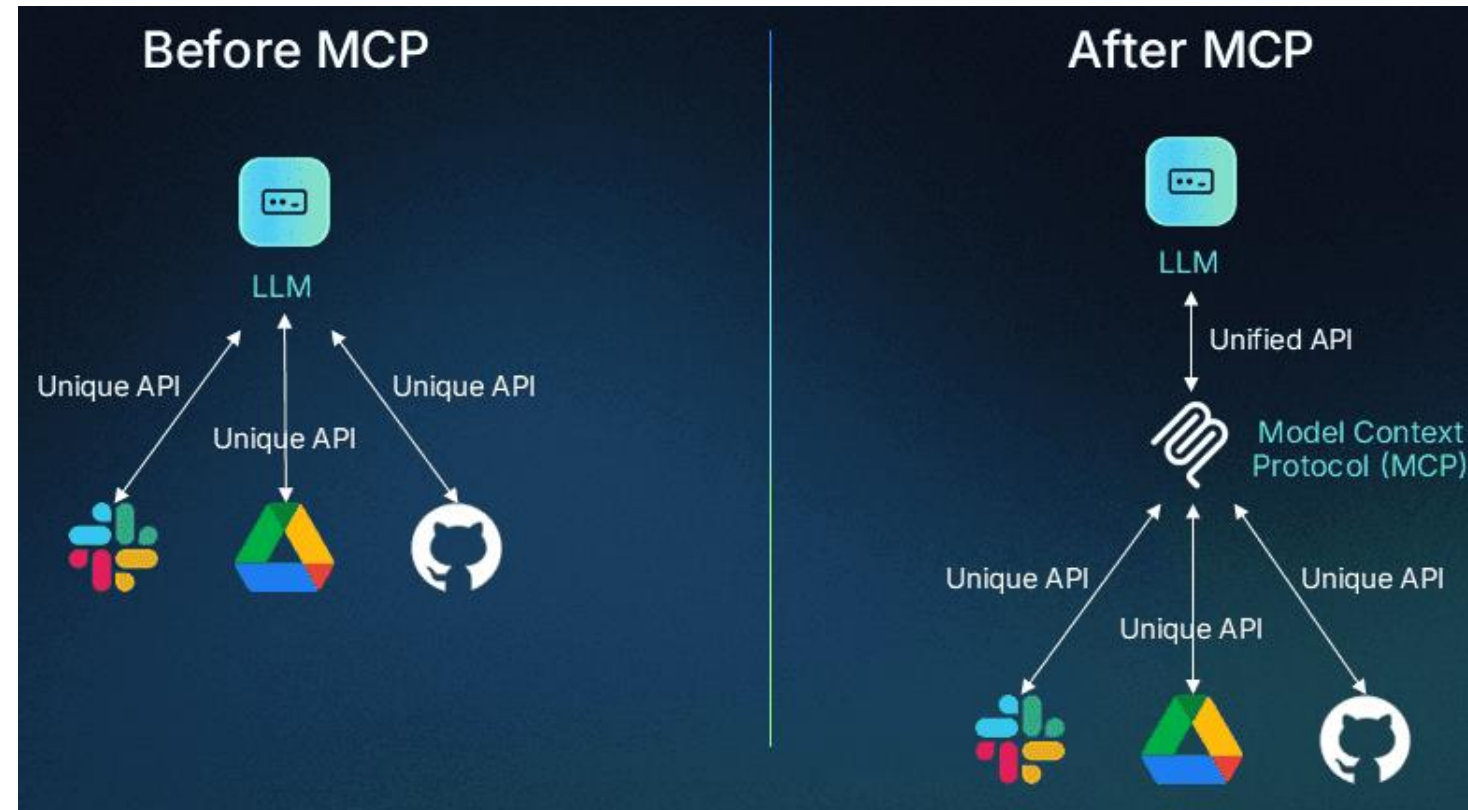
- Open standard for connecting AI applications to data sources and tools
- JSON-RPC based protocol enabling secure, controlled access to local and remote resources
- Developed by Anthropic, but provider-agnostic



Model Context Protocol (MCP)

- **Key Components:**

- **MCP Servers:** Expose resources (files, databases, APIs) and tools to AI applications
- **MCP Clients:** AI applications that consume MCP servers (Claude Desktop, IDEs, custom agents)
- **Transport Layer:** Local (stdio, pipes) or remote (HTTP, WebSocket) connections





MCP Impact On Agents

- Standardizes tool calling across different LLM providers
- Eliminates custom integration code for each data source
- Enables secure, permission-based access to sensitive resources
- Allows agents to discover available tools dynamically
- **Real Examples:**
 - **Filesystem MCP:** Read/write files with proper permissions
 - **Database MCP:** Query PostgreSQL, SQLite safely
 - **Git MCP:** Repository operations and code analysis
 - **Slack MCP:** Send messages, read channels
 - **Custom MCPs:** Company-specific APIs, internal tools



LLM-Powered Software Engineering Tools

- Examples of tools transforming dev workflows:
 - **Cursor**: AI-native code editor
 - **Claude Code**: Contextual understanding of codebases
 - Claude Code runs directly in your terminal
 - **Gemini Code Assist**: Task-specific completion and generation
 - **Lovable**: Zero-code chat interface for building software applications



LLM-Powered Software Engineering Tools

- What they're good at:
 - Generating new code from specs
 - Navigating, explaining, and refactoring large codebases
 - Accelerating test generation and debugging
- Caution: Not replacements for engineers — amplifiers of productivity



Challenges In LLM-Based Development

- Generates redundant or poorly integrated code
- Misses edge cases and dependencies
- Often fails to follow DRY principles
- May hallucinate unsafe practices
- Good engineering ≠ just good code snippets
- Building an application is easy, but building something scalable and secure is still challenging



Vibe Coding

- Vibe coding emphasizes staying in a creative flow:
 - The human developer avoids micromanaging the code, accepts AI-suggested completions liberally, and focuses more on iterative experimentation than code correctness or structure.
- The programmer shifts from manual coding to guiding, testing, and giving feedback about the AI-generated source code.
- Advocates of vibe coding say that it allows even amateur programmers to produce software without the extensive training and skills required for software engineering.
- Critics point out a lack of accountability and increased risk of introducing security vulnerabilities in the resulting software.



Project Implementation

Ideation Phase



Project Implementation

Development Phase



AGI

- What is AGI?
 - Hypothetical future system with general reasoning ability
 - The definition of AGI keeps changing
 - Goes beyond task-specific competence
 - Sam Altman: very tiny model with superhuman reasoning, 1 trillion tokens of context, and access to every tool you can imagine
- Why it's talked about:
 - Buzzword powering Big Tech investor hype
 - Timelines range wildly (1–20 years)
 - Often used aspirationally, not concretely



Contact Information

- Email:
 - maxmoundas@gmail.com
 - max.moundas@vanderbilt.edu
- LinkedIn: <https://www.linkedin.com/in/maxmoundas/>
- GitHub: <https://github.com/maxmoundas>